

Reproducing ReLA/GRES on gRefCOCO: failure-mode analysis and a post-hoc abstain-or-clarify detector

Visual Media (Eizo Media-gaku) — Final Report

Every number and figure in this report is generated by the scripts in `shiinayane/visual-media-gres` (`scripts/`, `results/`) and regenerates from one inference pass per split.

1. Identity

Name: WANG Yankai **Student ID:** 48256454

Department: 情報理工学系研究科 電子情報学専攻

Laboratory: 鈴木研究室 (Suzumura Lab)

Research topic: Vision–Language–Action / human–robot collaboration.

2. Target paper

GRES: Generalized Referring Expression Segmentation (Liu, Ding, Jiang, *CVPR 2023 Highlight*; network **ReLA**) generalizes classic Referring Expression Segmentation, which assumes *exactly one* target per expression. GRES additionally allows *multi-target* expressions (“the two bottles on the left”) and *no-target* expressions, whose referent is absent from the image. To support and measure this, the authors release **gRefCOCO** (COCO train2014 images, RefCOCO-style expressions; train/val/testA/testB splits) together with an evaluation protocol of five metrics:

- **gIoU** — mean per-sample IoU (a correctly predicted no-target sample scores 1);
- **cIoU** — cumulative intersection over cumulative union;
- **N-acc** — no-target recall, $TP/(TP + FN)$;
- **T-acc** — target retention, $TN/(TN + FP)$;
- **Pr@{0.7,0.8,0.9}** — precision at IoU thresholds 0.7/0.8/0.9.

ReLA is built on Mask2Former/Detectron2 and reaches state of the art on gRefCOCO.

3. Understanding of the method

Reading the released code (`gres_model/`), the inference path is:

1. A **backbone** (Swin-T/-B or R50) extracts multi-scale visual features and a **BERT-base** encoder embeds the expression; the backbone is language-aware (early fusion).
2. An **MSDeformAttn pixel decoder** builds a high-resolution mask feature map.
3. The **ReLA decoder** (`MultiScaleMaskedReferringDecoder`): 100 learnable *region queries* iterate through (i) Region–Image cross-attention, (ii) a first-layer Region–Language cross-attention gated by a learned `lang_weight`, and (iii) region–region self-attention. Two heads matter at test time: a *minimap/mask* branch yielding a 2-channel mask $\text{pred_masks} \in \mathbb{R}^{2 \times H \times W}$ (`argmax` over channels is the binary output), and a *no-target (NT)* head—a small MLP applied per region and **averaged over all 100 regions** to a single 2-way NT logit.
4. **Inference** sigmoids both outputs; the evaluator compares `argmax` of each to GT.

One observation motivates the rest of the report: the model already exposes two scalars at inference—the no-target score $\sigma(\text{nt_logit})$ and the foreground mask confidence—exactly the quantities needed to decide whether to act, abstain, or ask. ReLA collapses each with a

Table 1: Working software stack for the Blackwell/aarch64 + CUDA-13 server (single GPU, inference).

component	version	note
Python	3.11	venv
PyTorch	2.7.1+cu128	from the cu128 wheel index (the default PyPI aarch64 wheel is CPU-only)
torchvision	0.22.1	matched to torch 2.7.1
detectron2	0.6 (from source)	CPU-only C++ ops via <code>FORCE_CUDA=0</code> (avoids compiling CUDA kernels with the mismatched system nvcc 13.0)
numpy	1.26.4	pinned <2 for the pycocotools / detectron2 ABI
transformers	4.40.2	BERT-base text encoder

hard `argmax` at 0.5 and always commits to a mask or “empty.” Section 6 keeps the trained model unchanged and replaces this hard decision with a calibrated abstain-or-clarify policy.

4. Execution environment

Hardware. NVIDIA **GB10** (Grace-Blackwell, **aarch64**), CUDA driver 580 / **CUDA 13.0**, compute capability **sm_121**, 20 CPU cores, 119 GB RAM. The stack the paper targets (torch 1.11 / CUDA 11.8 / detectron2 0.6, Python 3.7–3.9) does not run on a Blackwell/aarch64 machine, so a modern equivalent was assembled (single GPU, inference only; Table 1).

Two source patches (`patches/`, applied by `setup.sh`), both forced by the Blackwell + CUDA-13 mismatch:

1. **MSDeformAttn fallback.** The hand-written CUDA op cannot be compiled (system nvcc 13.0 vs torch cu128/12.8); I fall back to ReLA’s own pure-PyTorch `ms_deform_attn_core_pytorch` (`F.grid_sample`), the reference implementation—numerically close, not bit-identical.
2. **sm_121 nVRTC fix.** `spatial_shapes.prod(1)` on int64 falls to the nVRTC iterator, whose cu128 build rejects `-arch=compute_121`; I replace it with the identical elementwise product `shapes[:,0]*shapes[:,1]` (precompiled kernel).

Baseline reproduction (gRefCOCO val, Swin-T, 14,229 expressions): gIoU = 55.76, cIoU = 55.64, N-acc = 46.3%, T-acc = 99.9%, Pr@0.7 = 66.5. The paper reports Swin-T cIoU 57.73 / gIoU 56.86, so the result lands within ~ 1 –2 points—a faithful reproduction; the small gap is consistent with the `grid_sample` fallback and single-GPU evaluation and does not affect any conclusion below. Inference runs at ≈ 0.24 s/sample; one pass dumps compact per-sample signals to `results/*.pkl` and all downstream analysis is offline.

5. Failure case analysis

I study two failure conditions on gRefCOCO’s official splits, so each claim is a measured rate, not an anecdote. F1 (no-target) is analysed on **val** (63% of its samples are no-target—the ideal stress test); F2 (multi-target) on **testA**, the split that contains a single-target baseline to compare against.

5.1 F1 — No-target hallucination (primary)

Setup. Restrict to the no-target subset of val (`gt_nt=True`, $n=8,905$, 63% of val). A *hallucination* is a no-target sample on which the model still emits a non-empty mask.

Result. The model hallucinates on **53.7%** of no-target samples (4,784/8,905); N-acc = 46.3%. The opposite error is almost absent: T-acc = 99.9% (only 7 of 5,324 present-target samples are wrongly called empty). The hallucinated masks are not small




Figure 1: No-target score for present vs absent referents. On absent referents the score is bimodal; the large low mode collapses inside the present-target distribution, so 53% fall below 0.5 and a mask is committed.




Figure 2: Each row: one *no-target* expression (left), the empty ground truth (middle), and the model’s confident hallucinated mask in red (right).

artefacts—median 7.2% of the image (mean 9.1%), mean foreground confidence 0.68: the model tends to be *wrong with high confidence*.

Mechanism. Figure 1 plots the model’s own no-target score $s_{nt} = \sigma(\text{nt_logit})_1$ for present vs absent referents. Present-target samples pile up at $s_{nt} \approx 0$ (mean 0.012). Absent-referent samples are **bimodal**: a high mode ($\sim 37\%$ above 0.9, correctly “no target”) and a **low mode** ($\sim 38\%$ collapsing to $s_{nt} \approx 0$, inside the present distribution), so overall 53% fall below the 0.5 threshold (the empty-subset mean 0.48 is the average of two modes, not a typical value). Architecturally the NT head is a *single* scalar **averaged over all 100 region queries**; a plausible account is that a few regions confidently matching a salient distractor pull the mean below 0.5 and a mask is committed. I logged only the pooled s_{nt} , not per-region logits, so this is consistent with the architecture but not directly verified (Section 8). For a system that acts on language, this high-confidence false positive is the most costly error.

5.2 F2 — Multi-target under-segmentation

Setup. I use **testA** because (unlike val, whose targeted expressions are all multi-instance) it has a single-instance baseline (5,917 one-target, 5,940 two-target, 2,895 three-or-more). For each GT instance I record $\text{coverage} = |\text{pred} \cap \text{inst}|/|\text{inst}|$, call it a *hit* if ≥ 0.5 , and report per-instance recall.

Result (two effects). (i) Per-instance recall is **flat from one to two targets**—0.868 (single) vs 0.881 (two)—then **collapses to 0.594 at ≥ 3 targets**. (ii) Independently of count, coverage within a multi-target expression is **very uneven**: the best-covered instance averages 0.92 but the worst only 0.58, and the model **drops at least one instance entirely in 37.8%** of multi-target testA samples (26.1% on val). The under-segmentation is thus not “two worse than one on average”; it is a **capacity ceiling at ≥ 3 targets** plus **lopsided coverage** sacrificing the less salient instance.

Mechanism. ReLA predicts a single 2-channel “target” mask via an einsum over per-region




Figure 3: testA. *Left:* per-instance recall is flat for 1–2 targets (0.87, 0.88) but collapses at ≥ 3 (0.59). *Right:* within a multi-target expression the best instance is covered at 0.92 but the worst at only 0.58 (38% of samples drop one entirely).




Figure 4: Qualitative multi-target cases. Each row is one expression; ground-truth instances in green (middle), prediction in red (right). The model frequently captures the salient instance and drops the other.

minimap weights. One mask can represent one region, or two when the soft weighting is shared, but cannot spread mass across many disjoint instances—matching both the sharp drop at ≥ 3 and the dropped weaker instance within a pair.

6. Improvement: a post-hoc abstain-or-clarify detector

Motivated by my research on human–robot collaboration—where a system facing an unsatisfiable or ambiguous instruction should ask or abstain rather than commit—I keep the trained model unchanged (no retraining) and replace its hard `argmax` with a calibrated policy. For each sample I define $s_{\text{abstain}} = \max(s_{\text{nt}}, 1 - \overline{\text{mask_conf}})$, high when the model leans “no target” or its committed mask is low-confidence. The policy is:

Table 2: gRefCOCO **val** ($n=14,229$). *In-sample:* τ is both chosen and scored here, so read this as the calibration curve, not a result.

op. point	gIoU	cIoU	N-acc	T-acc
baseline	55.76	55.64	46.3	99.9
safe 0.38	66.9	58.2	65.0	95.1
gIoU-opt 0.23	74.1	52.4	87.7	62.5

Table 3: gRefCOCO **testA** (held out, τ transferred unchanged; $n=19,200$). The honest evaluation.

op. point	gIoU	cIoU	N-acc	T-acc
baseline	65.03	65.42	50.6	99.0
safe 0.38	65.9	64.5	63.3	90.6
gIoU-opt 0.23	54.1	49.3	83.6	54.6




Figure 5: Abstain/clarify trade-off, calibrated on val. *Left:* N-acc vs T-acc as τ sweeps; the default **argmax** point (black) sits at the extreme no-abstention corner, the safe (orange) and gIoU-optimal (green) points move up the frontier. *Right:* gIoU and cIoU vs τ , both peaking above the default (dotted).

- $s_{\text{abstain}} \geq \tau \Rightarrow$ **ABSTAIN** (emit empty; scored as no-target);
- else if a mask has ≥ 2 connected components, the largest covering $< \rho$ of the area $\Rightarrow$ **CLARIFY** (“did you mean A or B?”);
- else $\Rightarrow$ **COMMIT** the mask.

τ is calibrated on val (max gIoU) and applied unchanged to held-out testA. I also report a safety point τ_{safe} : the most aggressive abstention keeping T-acc ≥ 0.95 . CLARIFY has no benchmark counterpart, so for scoring I keep the best-guess mask and report its trigger count separately. Calibration gives $\tau = 0.23$ (gIoU-optimal), $\tau_{\text{safe}} = 0.38$, $\rho = 0.7$.

7. Before/after comparison

τ is calibrated on **val** and applied **unchanged** to held-out **testA** (Tables 2–3; trade-off in Fig. 5).

On testA, with τ transferred unchanged, the detector is best described as a **safety re-weighting** rather than a segmentation-quality gain: the safe point raises N-acc by 12.7 points (50.6 $\rightarrow$ 63.3) at a cost of 8.4 points of T-acc (99.0 $\rightarrow$ 90.6), with overlap metrics essentially unchanged (gIoU +0.9, within the reproduction gap; cIoU -1.0). What it buys is a tunable, calibrated willingness to abstain on unsatisfiable instructions at a bounded retention cost. The aggressive gIoU-optimal τ *does not transfer*: tuned to val’s 63%-no-target prior, on the 23%-no-target testA it over-abstains and lowers gIoU by 11 points. Hence only the T-acc-constrained point transfers safely, and a practical detector would calibrate τ to the deployment base-rate. The **CLARIFY** branch (separate from abstention) fired on 1,131 val / 1,712 testA committed predictions at $\tau=0.23$ —multi-component masks for which “which one?” is a reasonable alternative; an illustrative remedy for F2, reported as a trigger count since gRefCOCO has no clarification label.

8. Limitations

- Inference-only on the frozen Swin-T checkpoint (Swin-B downloaded but, per the resource guideline, not evaluated); with no retraining the detector reuses signals, not fixes the F2 bias.
- The `grid_sample` MSDeformAttn fallback is the reference op (close, not identical); the ~ 1 -point baseline gap is plausibly but not provably benign.
- The F1 mechanism is a hypothesis: I logged only the pooled s_{nt} , not per-region NT logits, so the averaging-dilution account is consistent with, but not proof of, the architecture.
- $\max(s_{nt}, 1 - \text{mask_conf})$ is a simple interpretable rule with correlated terms; a learned calibrator over $\{s_{nt}, \text{mask_conf}, \text{area}, \#\text{comp}\}$ would likely do better.
- On held-out data the detector is a safety re-weighting, not a quality gain, and a single global τ depends on the no-target base-rate. The CLARIFY trigger is a geometric heuristic reported only as a count (no clarification label; it does not yet read the expression).

9. Generative-AI usage

I used a generative-AI coding assistant (Anthropic’s Claude, via the Claude Code CLI) throughout, and disclose it as follows. **The AI assisted with:** debugging the Blackwell/CUDA-13 environment and proposing the working stack and the two source patches (Section 4); writing the bulk of the analysis code in `scripts/` (inference loop, metric port, failure-mode and detector scripts, plotting); running the experiments; and drafting report text from the resulting numbers. **I am responsible for:** selecting the paper and the direction (the abstain-or-clarify framing tied to my HRC research); reviewing the code and setup; checking every reported number and figure against the generated tables; requesting the corrections from an internal review (e.g. re-centring F2 on testA; reporting held-out rather than in-sample detector results); and editing the report into final form. A detailed running log is in `GENAI_USAGE_LOG.md`.

10. Discussion

The two failure modes share an origin: ReLA makes a single hard decision at inference (one mask, one 0.5 no-target threshold), whereas the generalized task is partly about *uncertainty*—whether a referent exists at all, and how many. The no-target and multi-target subsets are exactly where committing to a single `argmax` is the wrong default. The detector adds no capacity; it reuses the uncertainty the model already encodes and turns it into three actions (act, abstain, ask), giving a calibrated, tunable abstention at a bounded retention cost that transfers to a held-out split. For a benchmark aimed at human–robot use, this argues for an explicit abstain/clarify option in how grounding models are both *scored* and *operated*, since these are the cases where a confident-but-wrong mask is most costly. The gains are modest and prior-dependent; natural next steps are a small learned calibrator over the four signals, an expression-aware CLARIFY trigger, and an abstention-aware metric (e.g. risk–coverage) reported alongside gIoU/cIoU.

Citation: Chang Liu, Henghui Ding, Xudong Jiang. “GRES: Generalized Referring Expression Segmentation.” **CVPR 2023** (Highlight). arXiv:2306.00968.